

1 **Problem 1 (Extended Kalman Filter, 25 points).** In this problem, we will see how  
 2 to use filtering to estimate an unknown system parameter. Consider a dynamical  
 3 system given by

$$\begin{aligned} x_{k+1} &= ax_k + \epsilon_k \\ y_k &= \sqrt{x_k^2 + 1} + \nu_k \end{aligned} \quad (1)$$

4 where  $x_k, y_k \in \mathbb{R}$  are scalars,  $\epsilon_k \sim N(0, 1)$  and  $\nu_k \sim N(0, 1/2)$  are zero-mean  
 5 scalar Gaussian noise uncorrelated across time  $k$ . The constant  $a$  is unknown and  
 6 we would like to estimate its value. If we know that our initial state has mean 1 and  
 7 variance 2

$$x_0 \sim N(1, 2),$$

8 develop the equations for an Extended Kalman Filter (EKF) to estimate the unknown  
 9 constant  $a$ .

10 (a) **(5 points)** You should first simulate (1) with  $a = -1$ . This is the ground-  
 11 truth value of  $a$  that we would like to estimate. Provide details of how you  
 12 simulated the system, in particular how did you sample the noise  $\epsilon_k, \nu_k$ .  
 13 The observations  $D = \{y_k : k = 1, \dots\}$  are the “dataset” that we thus  
 14 collect from the system. Run the simulation for about 100 observations.

15 (b) **(15 points)** You should now develop the EKF equations that will use the  
 16 collected dataset  $D$  to estimate the constant  $a$ . Discuss your approach in  
 17 detail. Your goal is to compute two quantities

$$\begin{aligned} \mu_k &= \mathbb{E}[a_k | y_1, \dots, y_k] \\ \sigma_k^2 &= \text{var}(a_k | y_1, \dots, y_k). \end{aligned}$$

18 for all times  $k$ .

19 (c) **(5 points)** Plot the true value  $a = -1$ , and the estimated values  $\mu_k \pm \sigma_k$  as  
 20 a function of time  $k$ . Discuss your result. In particular, do your estimated  
 21 values  $\mu_k \pm \sigma_k$  match the ground-truth value  $a = -1$ ? Does the error  
 22 reduce as your incorporate more and more observations?

26 **Problem 2 (Unscented Kalman Filter, 100 points).** In this problem, you will imple-  
 27 ment an Unscented Kalman Filter (UKF) to track the orientation of a robot in three-  
 28 dimensions. We have given you observations from an inertial measurement unit  
 29 (IMU) that consists of gyroscopes and accelerometers and corresponding data from a  
 30 motion-capture system (called “Vicon”, see <https://www.youtube.com/watch?v=qgS1pwsHQIA>  
 31 for example). We will develop the UKF for the IMU data and the vicon data for  
 32 calibration and tuning of the filter, this is typical of real applications where the robot  
 33 uses an IMU but the filter running on the robot will be calibrating before test-time  
 34 in the lab using an expensive and accurate sensor like a Vicon.

1 (a) **Loading and understanding the data:** First, load the data given on Canvas  
2 (file “hw2\_p2\_data.zip”) using code of the form.

```
3 from scipy import io
4
5
6 data_num = 1
7 imu = io.loadmat('imu/imuRaw'+str(data_num)+'.mat')
8 accel = imu['vals'][0:3,:]
9 gyro = imu['vals'][3:6,:]
10 T = np.shape(imu['ts'])[1]
```

12 Ignore other fields inside the .mat file, we will not use them.

13 You can use the following code to load the vicon data

```
14
15 vicon = io.loadmat('vicon/viconRot'+str(data_num)+'.mat')
```

17 while calibrating and debugging. But do not include this line in the autograder  
18 submission because we do not store the vicon data on the server.

19 (b) **(15 points) Calibrating the sensors.** Check the arrays accel and gyro. The  
20 former gives the observations received by the accelerometer inside the IMU and  
21 the latter gives observations from gyroscope. The variable  $T$  denotes the total  
22 number of time-steps in our dataset. You will have to read the IMU reference  
23 manual (file “imu\_reference.pdf” on Canvas) to understand the quantities stored  
24 in these arrays. Pay careful attention to the following things. First, the accel/gyro  
25 readings are integers and not metric quantities, this is because there is usually an  
26 analog-to-digital conversion (ADC) that happens in these sensors and one reads off  
27 the ADC value as the actual observation. Because of the way these MEMS sensors  
28 are constructed, they will have biases and sensitivity with respect to the working  
29 voltage. In order to convert from raw values to physical units, the equation for both  
30 accel and gyro is typically

$$\text{value} = (\text{raw} - \beta) \frac{3300 \text{ mV}}{1023 \alpha}$$

31 where  $\beta$  called the bias, mV stands for milli-volt (most onboard electronics operators  
32 at 3300 mV) and  $\alpha$  is the sensitivity of the sensor. For the accelerometer,  $\alpha$  has  
33 units of mV/g where g refers to the gravitational acceleration  $9.81 \text{ m/s}^2$ . In other  
34 words, if  $\alpha = 100 \text{ mV/g}$  and bias  $\beta$  is zero, and if the raw accelerometer reading is  
35 10, the actual value of the acceleration along that axis is

$$\text{value} = 10 \times \frac{3300}{1023 \times 100} \times 9.81 = 3.16 \text{ m/s}^2$$

36 Similarly, the sensitivity of a gyroscope has units mV/degrees/sec. You will have to  
37 convert the sensitivity into mV/radians/sec to make everything into consistent units.

38 Typically, in a real application, we do not know the bias and sensitivity of either  
39 sensor. Your goal is to use the rotation matrices in the Vicon data as the ground-truth  
40 orientation (see section on quaternions below) to estimate the bias and sensitivity of  
41 both the accelerometer and the gyroscope. You should be careful on two counts.

- (1) The orientation of the IMU need not be the same as the orientation of the Vicon coordinate frame. Plot all quantities in the arrays `accel`, `gyro` and `vicon` rotation matrices to make sure you get this right. Do not proceed to implementing the filter if you are not convinced your solution for this part is correct.
- (2) The acceleration  $a_x$  and  $a_y$  is flipped in sign due to device design. A positive acceleration in body-frame will result in a negative number reported by the IMU. See the IMU manual for more insight.

*Hint:* To find the sensitivity for the accelerometer, we can assume the only force acting is the gravitational force. Then the magnitude of your 3-dimensional accelerometer readings should be as close to 9.81 as possible. Plot the roll, pitch, and yaw values from the Vicon data; you can extract these from the Vicon rotation matrix. You should compare the Vicon plots with some simple plots obtained only from the accelerometer, and separately only the gyroscope values, to predict orientation. From the accelerometer, you can directly compute roll and pitch for each timestep and compare these with ground truth (Vicon) data to ensure your sensitivity and axes are correct. For the gyroscope, you can use the initial orientation from the accelerometer and then integrate the angular velocity values from the gyroscope for the rest of the time series. Note that the orientation estimates you get from this method will have significant drift, but you should be able to get a sense of the scale and check your sensitivity values. The purpose here is to ensure you are converting the raw digital values in the dataset into meaningful physical units before we begin the filtering.

(c) **(0 points) Quaternions for orientation** We have given you a file named `quaternion.py` that implements a Python class for handling quaternions. Read this code carefully. In particular, you should study the function `euler_angle` which returns the Euler angles corresponding to a quaternion, `from_rotm` which takes in a  $3 \times 3$  rotation matrix and assigns the quaternion and the function `__mul__` which multiplies two quaternions together. Try a few test cases for converting to-and-fro from a rotation matrix/Euler angles to a quaternion to solidify your understanding here.

(d) **(0 points) Implementing the UKF** Given this setup, you should now read the PDF on Canvas titled “hw2\_p2\_ukf\_writeup” to implement an Unscented Kalman Filter for tracking the orientation using these observations. The state of your filter will be


$$x = \begin{bmatrix} q \\ \omega \end{bmatrix} \in \mathbb{R}^7$$

where  $q$  is the quaternion that indicates the orientation and  $\omega$  is the angular velocity. The observations are of course the readings of the accelerometer and the gyroscope that we discussed above; recall that gyroscopes measure the angular velocity  $\omega$  itself

1 which simplifies their observation model. The paper also has a magnetometer as a  
2 sensor (which measures the orientation with respect to the magnetic north pole) but  
3 we will not use it here. You should implement quaternion averaging as described in  
4 Section 3.4 of the paper; this is essential for the UKF to work. **You will have to**  
5 **choose yourself the values of the initial covariance of the state, dynamics noise**  
6 **and measurement noise.** You can discuss your steps and choices in your solution  
7 PDF if you want us to follow through your implementation.

8 (e) **(10 points) Analysis and debugging** Plot the quaternion  $q$  (mean and di-  
9 agonal of covariance), the angular velocity  $\omega$  (mean and diagonal of covariance),  
10 the gyroscope readings in rad/sec and the quaternion corresponding to the vicon  
11 orientation as a function of time in your solution PDF. Do not plot on the server, it  
12 may crash out. You should show the results for one dataset and discuss whether your  
13 filter is working well. You should also use these plots to debug your performance  
14 on the other datasets; plotting everything carefully is the fastest way to debugging  
15 the UKF.

12 **Problem 3 (Simultaneous Localization and Mapping (SLAM) with a particle**  
13 **filter, 80 points. Do not use Google Colab to do this homework).** In this problem,  
14 we will implement mapping and localization in an indoor environment using infor-  
15 mation from an IMU and a LiDAR sensor. We have provided you data collected  
16 from a humanoid named THOR that was built at Penn and UCLA  
17 (<https://archive.darpa.mil/roboticschallenge/finalist/thor.html>). You can read more  
18 about the hardware in this paper (<https://ieeexplore.ieee.org/document/7057369>.)



19

20 **Hardware setup of Thor** The humanoid has a Hokuyo LiDAR sensor ([https://hokuyo-](https://hokuyo-usa.com/products/lidar-obstacle-detection)  
21 [usa.com/products/lidar-obstacle-detection](https://hokuyo-usa.com/products/lidar-obstacle-detection) on its head (the final version of the robot  
22 had it in its chest but this is a different version); details of this are in the code (which  
23 will be explained shortly). This LiDAR is a planar LiDAR sensor and returns 1080  
24 readings at each instant, each reading being the distance of some physical object  
25 along a ray that shoots off at an angle between  $(-135, 135)$  degrees with discretiza-  
26 tion of 0.25 degrees in an horizontal plane (shown as white rays in the picture).

1 We will use the position and orientation of the head of the robot to calculate the  
2 orientation of the LiDAR in the body frame.

3 The second kind of observations we will use pertain to the location of the robot.  
4 However, in contrast to the previous homework where we used the raw accelerometer  
5 and gyroscope readings to get the orientation, we will directly use the  $(x, y, \theta)$  pose  
6 of the robot in the world coordinates ( $\theta$  denotes yaw). These poses were created  
7 presumably on the robot by running a filter on the IMU data (such estimates are  
8 called odometry estimates), and just as you saw some tracking errors in the previous  
9 homework, these poses will not be extremely accurate. However, we will treat them  
10 conceptually the same way as we treated Vicon in the previous homework, namely  
11 as a much more precise estimate of the pose of the robot that is used to check how  
12 well SLAM is working.

13 **Coordinate frames** The body frame is at the top of the head (X axis pointing  
14 forwards, Y axis pointing left and Z axis pointing upwards), the top of the head is at  
15 a height of 1.263m from the ground. The transformation from the body frame to the  
16 LiDAR frame depends upon the angle of the head (pitch) and the angle of the neck  
17 (yaw) and the height of the LiDAR above the head (which is 0.15m). The world  
18 coordinate frame where we want to build the map has its origin on the ground plane,  
19 i.e., the origin of the body frame is at a height of 1.263m with respect to the world  
20 frame at location  $(x, y, \theta)$ .

## 21 **Data and code**

22 (a) **(0 points)** We have provided you 4 datasets corresponding to 4 different  
23 trajectories of the robot in Towne Building at Penn. For example, dataset 0 consists  
24 of two files `data/train/train_lidar0.mat` and `data/train/train_joint0.mat` which contain  
25 the LiDAR readings and joint angles respectively. The functions `load_lidar_data`  
26 and `load_joint_data` inside `load_data.py` read the data. You can run the function  
27 `show_lidar` to see the LiDAR data. Each of the data reading functions returns  
28 a data-structure where  $t$  refers to the time-stamp (in seconds) of the data, `xyth`  
29 refers to  $(x, y, \theta)$  *pose of the LiDAR* and `rpy` refers to Euler angles (roll, pitch,  
30 yaw). The joint data contains a number of fields, but we are only interested in the  
31 angle of the head and the neck at a particular time-stamp. You should read these  
32 functions carefully and check the values returned by them. The dicts `joint_names`  
33 and `joint_names_to_index` can be used to read off the data of a specific joint (we  
34 only need the head and the neck).

35 (b) **(0 points)** Next look at the `slam.py` file provided to you. Read the code  
36 for the class `map_t` and `slam_t` and the comments provided in the code very  
37 carefully. You are in charge of filling in the missing pieces marked as `TODO`:  
38 `XXXXXX`. A suggested order for studying this code is as follows: `slam_t.read_data`,  
39 `slam_t.init_sensor_model`, `slam_t.init_particles`, `slam_t.rays2world`, `map_t.__init__`,  
40 `map_t.grid_cell_from_xy`. Next, the file `utils.py` contains a few standard rigid-body  
41 transformations that you will need. You should pay attention to the functions

1 smart\_plus\_2d and smart\_minus\_2d that will be used to code up the dynamics  
2 propagation step of the particle filter.

3 (c) **(20 points, dynamics step)** Next look at main.py which has two functions  
4 run\_dynamics\_step and run\_observation\_step which act as test functions to check  
5 if the particle filter and occupancy grid update has been updated correctly. The  
6 run\_dynamics function plots the trajectory of the robot (as given by its IMU data  
7 in the LiDAR data-structure). It also initializes 3 particles and plots all particles  
8 at different time-steps while performing the dynamics step with a very small dy-  
9 namics noise; this is a very neat way of checking if dynamics propagation in the  
10 particle filter is working correctly. This function will create two plots, one for the  
11 odometry trajectory and one more for the particle trajectories, both these trajec-  
12 tories should match after you code up the dynamics function slam\_t.dynamics\_step  
13 correctly. Include these plots for all datasets in your report. Briefly explain how you  
14 implemented the dynamics step.

15 (d) **(20 points, observation step)** The function run\_observation\_step is used to  
16 perform the observation step of the particle filter to get an estimate of the location  
17 of the robot and updates to the occupancy grid using observations from the LiDAR.  
18 First read the comments for the function slam\_t.observation\_step carefully.

19 We first discuss the particle filtering part.

20 (i) Compute the head and neck position for the time  $t$ . For each particle,  
21 assuming that that particle is indeed the true position of the robot, project  
22 the LiDAR scan `slam_t.lidar[t]['scan']` into the world coordinates using the  
23 `slam_t.ray2world` function. The end points of each ray tell us which cells in  
24 the map are occupied, for each particle.

25 (ii) In order to compute the updated weights of the particle, we need to know  
26 the likelihood of LiDAR scans given the state (our current occupancy grid  
27 in the case of SLAM). We are going to use a simple model to do so

$$\log P(\text{LiDAR scan as if the robot is at particle } p \mid m) = \sum_{ij \in O} m_{ij} \quad (1)$$

28 where  $O$  is the set of occupied cells as detected by the LiDAR scan assuming  
29 the robot is at particle  $p$  and  $m_{ij}$  is our current estimate of the binarized  
30 map (more on this below). In simple words, if the occupied cells as given  
31 by our LiDAR match the occupied cells in the binarized map created from  
32 the past observations, then we say the log-probability of particle  $p$  is large.

33 (iii) You will next implement the function `slam_t.update_weights` that takes  
34 the log-probability of each particle  $p$ , its previous weights, calculates the  
35 updated weights of the particles.

36 (iv) Typically, resampling step (`slam_t.stratified_resampling`) is performed only  
37 if the effective number of particles (as computed in `slam_t.resample_particles`)  
38 falls below a certain threshold (30% in the code). Implement resampling as  
39 we discussed in the lecture notes.

1 We will now do the mapping part. We have a number of particles  $p^i = (x^i, y^i, \theta^i)$   
2 that together give an estimate of the distribution of the location of the robot. For  
3 this homework, you will only use the particle with the largest weight to update  
4 the map although typically we update the map using all particles. Our goal is  
5 simple: we want to increase `map_t.log_odds` array at cells that are recorded as  
6 obstacles by the LiDAR and decrease the values in all other cells. You should  
7 add `slam_t.log_odds_occ` to all occupied cells and add `slam_t.log_odds_free` from  
8 all cells in the map. It is also a good idea to clip the `log_odds` to like between  
9 `[-slam_t.map.log_odds_max, slam_t.map.log_odds_max]` to prevent increasingly  
10 large values in the `log_odds` array. The array `slam_t.map.cells` is a binarized version  
11 of the map (which is used above to calculate the observation likelihood).

12 Check the `run_observation_step` function after you have implemented the obser-  
13 vation step. Since the map is initialized to zero at the beginning of SLAM which  
14 results in all observation log-likelihoods to be zero in (1), we need to do something  
15 special for the first step. We will use the first entry in `slam_t.lidar[0]['xyth']` to  
16 get an accurate pose for the robot and use its corresponding LiDAR readings to  
17 initialize the occupancy grid. You can do this easily by initializing the particle filter  
18 to have just one particle and simply calling the `slam_t.observation_step` as shown in  
19 `main.py`.

20 Include in your report the output of the `run_observation_step` function for one  
21 time-step. Briefly explain how you implemented the observation step.

22 (e) **(40 points)** You will now run the full SLAM algorithm that performs one  
23 dynamics step and observation step at each iteration in the function `run_slam` in  
24 `main.py`. Make sure to start SLAM only after the time when you have both LiDAR  
25 scans and joint readings (the two arrays start at different times). For all 4 datasets,  
26 include in your report the plots of the final binarized version of the map, the  $(x, y)$   
27 location of the particle in the particle filter with the largest weight at each time-step  
28 and the odometry trajectory  $(x, y)$  (in a different color); this counts for 10 points  
29 each.